

Task Assignment: Ramayana Source Extraction, Translation & Structuring for RAG

Objective

Your task is to extract, clean, translate, and structure the content of the attached document named **ramayana.pdf** from Telugu into highly accurate, clean English source material — organized as retrieval-ready **chunks** — that can be best ingested by a Retrieval-Augmented Generation (RAG) system.

Accuracy is critical for this task. You are fully permitted to leverage any **DeepSeek** model (such as DeepSeek-Flash or DeepSeek-Pro) to ensure maximum translation fidelity, context retention, and formatting precision.

You are also free to use any coding tools or AI assistants of your choice (e.g., **Claude Code**, **Codex**, **GitHub Copilot**) to build, run, and refine the extraction pipeline.

Constraint: For the actual generation and translation of the content, use **DeepSeek models only** — do not use any other LLM APIs (e.g., Gemini, OpenAI / GPT) for content generation. Coding assistants such as Claude Code, Codex, or Copilot may still be used to build the pipeline itself.

Detailed Task Instructions

1. Document Reference & Scope

- Analyze the provided file **ramayana.pdf**.
- Process the text sequentially, capturing the metadata (preface/index if available) and moving systematically through the chapters (e.g., Bala Kanda Chapter 1, Chapter 2, etc.).

2. Data Cleaning & Text Normalization

- Strip out any OCR noise, typos, stray symbols, or printing artifacts.
- Standardize the spelling of historical/mythological names, locations, and Vedic terms (e.g., Nakshatras, Yagas, Kandas) according to standard English transliteration conventions.

3. High-Quality Translation (Telugu → English)

- Translate the entire text into clear, grammatically flawless English.
- Maintain the true cultural and theological context of the epic. Avoid overly modern paraphrasing that strips away the text's original meaning.

Goal & Output

The end goal is **chunks** — clean, faithful English source material broken into retrieval-ready chunks (each a coherent, self-contained passage) that a **Retrieval-Augmented Generation (RAG)** system can ingest effectively.

We are intentionally **not** prescribing a fixed output schema. You are free to take your own decisions on the chunking strategy — chunk size, boundaries, and any accompanying metadata — so that the chunks are best suited for high-quality RAG ingestion and retrieval.

Submission

Deadline: Submit within **24 hours** from the time the assignment is shared.

- Submit your work as a **GitHub repository** containing both the extraction/translation pipeline code and the final output (the cleaned, translated source material as RAG-ready **chunks**).
- Include a short **README** explaining how to run the pipeline and how the output is organized.
- Share the completed assignment (GitHub repository link) with **dolaakhildatta@gmail.com** and **neelalohithpenmetsa@gmail.com**.

Security — never commit secrets: Do not push any API keys, tokens, passwords, or credentials to the repository — **irrespective of whether the repo is private or public**. Keep secrets in environment variables or a `.env` file excluded via `.gitignore`, and follow standard secret-hygiene practices (e.g., review the diff before every commit).

Note: Feel free to recheck the quality of the output and accordingly refine the pipeline. For any questions about the assignment, feel free to reach out to us by email at the addresses above.